

# Druid pour l'analyse de données en temps réel

Yann Esposito

7 Avril 2016

## Abstract

Druid expliqué rapidement, pourquoi, comment.

## Contents

<b>1</b>	<b>Intro</b>	<b>3</b>
1.1	Plan . . . . .	3
1.2	Experience . . . . .	3
1.3	Real Time? . . . . .	3
1.4	Demande . . . . .	3
1.5	En pratique . . . . .	3
1.6	Origine (PHP) . . . . .	4
1.7	Introduction . . . . .	4
1.8	Pre Considerations . . . . .	4
1.9	Try to conceptualize (events) . . . . .	4
1.10	In the End . . . . .	5
<b>2</b>	<b>Druid</b>	<b>5</b>
2.1	Who . . . . .	5
2.2	Goal . . . . .	5
2.3	Concepts . . . . .	5
2.4	Features . . . . .	5
2.5	Storage . . . . .	5
2.6	Columnar Storage . . . . .	6

2.7	Index . . . . .	6
2.8	Data Segments . . . . .	6
2.9	Real-time ingestion . . . . .	6
2.10	Batch Ingestion . . . . .	6
2.11	Real-time Ingestion . . . . .	7
<b>3</b>	<b>Querying</b>	<b>7</b>
3.1	Query types . . . . .	7
3.2	Tip . . . . .	7
3.3	Query Spec . . . . .	7
3.4	Example(s) . . . . .	7
3.5	Caching . . . . .	8
3.6	Load Rules . . . . .	8
<b>4</b>	<b>Components</b>	<b>8</b>
4.1	Druid Components . . . . .	8
4.2	Coordinator . . . . .	8
4.3	Real-time Nodes . . . . .	9
4.4	Historical Nodes . . . . .	9
4.5	Overlord . . . . .	9
4.6	Middle Manager . . . . .	9
4.7	Broker Nodes . . . . .	9
4.8	Deep Storage . . . . .	9
<b>5</b>	<b>Considerations &amp; Tools</b>	<b>9</b>
5.1	When <i>not</i> to choose Druid . . . . .	9
5.2	Graphite (metrics) . . . . .	10
5.3	Pivot (exploring data) . . . . .	10
5.4	Caravel (exploring data) . . . . .	11

# **1 Intro**

## **1.1 Plan**

- Introduction; why?
- How?

## **1.2 Experience**

- Real Time Social Media Analytics

## **1.3 Real Time?**

- Ingestion Latency: seconds
- Query Latency: seconds

## **1.4 Demande**

- Twitter: 20k msg/s, 1msg = 10ko pendant 24h
- Facebook public: 1000 à 2000 msg/s en continu

## **1.5 En pratique**

- Twitter: 400 msg/s en continu, pics à 1500

## 1.6 Origine (PHP)



## 1.7 Introduction

- Traitement de donnée gros volume + faible latence
- Typiquement pulse

DEMO

## 1.8 Pre Considerations

Discovered vs Invented

## 1.9 Try to conceptualize (events)

Scalable + Real Time + Fail safe

- timeseries
- alerting system
- top N
- etc...

## 1.10 In the End

Druid concepts are always emerging naturally

# 2 Druid

## 2.1 Who

Metamarkets

## 2.2 Goal

Druid is an open source store designed for real-time exploratory analytics on large data sets.

hosted dashboard that would allow users to arbitrarily explore and visualize event streams.

## 2.3 Concepts

- Column-oriented storage layout
- distributed, shared-nothing architecture
- advanced indexing structure

## 2.4 Features

- fast aggregations
- flexible filters
- low latency data ingestion

**arbitrary exploration of billion-row tables tables with sub-second latencies**

## 2.5 Storage

- Columnar
- Inverted Index
- Immutable Segments

## 2.6 Columnar Storage

## 2.7 Index

- Values are dictionary encoded

`{"USA" 1, "Canada" 2, "Mexico" 3, ...}`

- Bitmap for every dimension value (used by filters)

`"USA" -> [0 1 0 0 1 1 0 0 0]`

- Column values (used by aggregation queries)

`[2,1,3,15,1,1,2,8,7]`

## 2.8 Data Segments

- Per time interval
- skip segments when querying
- Immutable
- Cache friendly
- No locking
- Versioned
- No locking
- Read-write concurrency

## 2.9 Real-time ingestion

- Via Real-Time Node and Firehose
- No redundancy or HA, thus not recommended
- Via Indexing Service and Tranquility API
- Core API
- Integration with Streaming Frameworks
- HTTP Server
- **Kafka Consumer**

## 2.10 Batch Ingestion

- File based (HDFS, S3, ...)

## 2.11 Real-time Ingestion

Task 1: [ Interval ] [ Window ]  
Task 2: [ ]  
----->  
time

Minimum indexing slots =  
Data Sources  $\times$  Partitions  $\times$  Replicas  $\times$  2

## 3 Querying

### 3.1 Query types

- Group by: group by multiple dimensions
- Top N: like grouping by a single dimension
- Timeseries: without grouping over dimensions
- Search: Dimensions lookup
- Time Boundary: Find available data timeframe
- Metadata queries

### 3.2 Tip

- Prefer `topN` over `groupBy`
- Prefer `timeseries` over `topN`
- Use limits (and priorities)

### 3.3 Query Spec

- Data source
- Dimensions
- Interval
- Filters
- Aggregations
- Post Aggregations
- Granularity
- Context (query configuration)
- Limit

### 3.4 Example(s)

TODO

### 3.5 Caching

- Historical node level
- By segment
- Broker Level
- By segment and query
- `groupBy` is disabled on purpose!
- By default - local caching

### 3.6 Load Rules

- Can be defined
- What can be set

## 4 Components

### 4.1 Druid Components

- Real-time Nodes
- Historical Nodes
- Broker Nodes
- Coordinator
- For indexing:
- Overlord
- Middle Manager
- Deep Storage
- Metadata Storage
- Load Balancer
- Cache

### 4.2 Coordinator

Manage Segments



### 4.3 Real-time Nodes

- Pulling data in real-time
- Indexing it

### 4.4 Historical Nodes

- Keep historical segments

### 4.5 Overlord

- Accepts tasks and distributes them to middle manager

### 4.6 Middle Manager

- Execute submitted tasks via Peons

### 4.7 Broker Nodes

- Route query to Real-time and Historical nodes
- Merge results

### 4.8 Deep Storage

- Segments backup (HDFS, S3, ...)

## 5 Considerations & Tools

### 5.1 When *not* to choose Druid


- Data is not time-series
- Cardinality is *very* high
- Number of dimensions is high
- Setup cost must be avoided

## 5.2 Graphite (metrics)



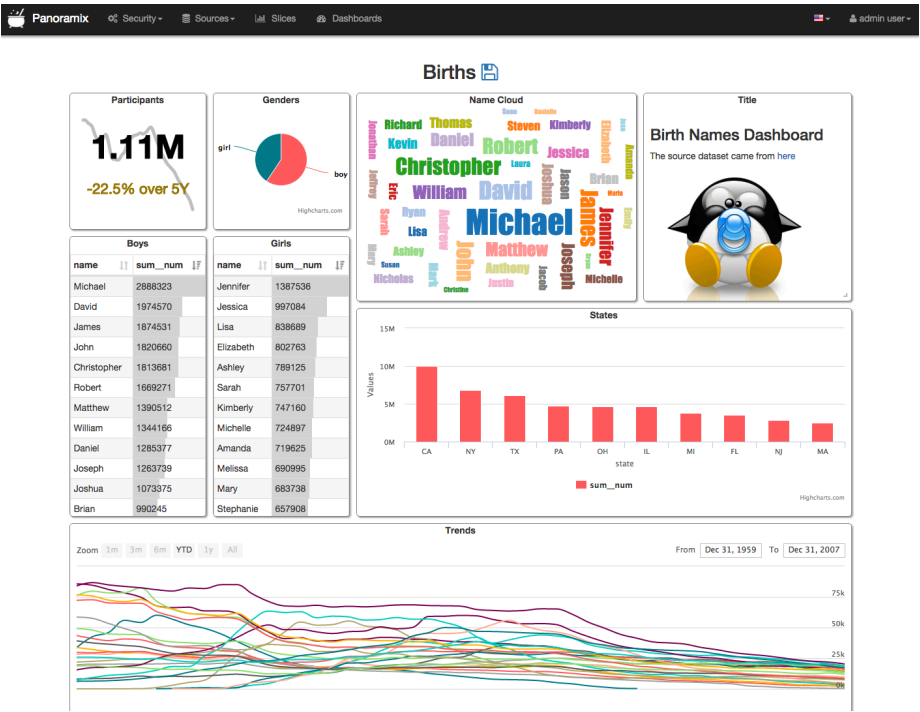
## Graphite

### 5.3 Pivot (exploring data)



## Pivot

5.4 Caravel (exploring data)



Caravel